Table 1 | Performance comparison on MPDocBench. All models in the table are end-to-end VLM-based architectures.

| Method | Overall↑ | Text$^{Edit}$ ↓ | T-Text$^{Edit}$ ↓ | Formula$^{CDM}$ ↑ | Table$^{TEDS}$ ↑ | T-Table$^{TEDS}$ ↑ | Figure$^{F1}$ ↑ | ReadOrder$^{Edit}$ ↓ | Heading$^{TEDS}$ ↑ |
|---|---|---|---|---|---|---|---|---|---|
| *General VLMs* | | | | | | | | | |
| Gemini-3.1-Pro Preview [9] | 71.79 | 0.070 | 0.232 | 88.37 | 81.68 | 61.26 | 59.08 | 0.127 | 26.94 |
| Qwen3.8-Max [25] | 78.30 | 0.057 | 0.196 | 89.24 | 85.89 | 67.18 | 70.19 | 0.097 | 48.80 |
| Kimi-K3 [24] | 60.77 | 0.162 | 0.551 | 41.16 | 75.45 | 81.57 | 40.47 | 0.206 | 39.30 |
| Qwen3.6-plus [25] | 72.14 | 0.096 | 0.245 | 88.77 | 82.99 | 60.50 | 64.80 | 0.179 | 32.06 |
| Qwen3-VL-235B [1] | 73.98 | 0.087 | 0.187 | 84.71 | 81.35 | 61.25 | 63.25 | 0.136 | 42.34 |
| InternVL-3.5-38B [27] | 57.01 | 0.133 | 0.515 | 84.30 | 69.86 | 51.69 | 7.97 | 0.199 | 27.01 |
| Minimax-M3 [16] | 62.65 | 0.167 | 0.422 | 81.30 | 77.91 | 76.51 | 7.18 | 0.231 | 40.39 |
| Gemini-3.8-Flash [10] | 76.25 | 0.081 | 0.122 | 83.27 | 83.15 | 87.41 | 62.86 | 0.294 | 43.05 |
| Qwen-3.8-Flash-Next [25] | 77.36 | 0.078 | 0.163 | 86.70 | 86.24 | 88.36 | 64.31 | 0.277 | 45.09 |
| DeepSeek-4.1-Flash [32] | 66.25 | 0.114 | 0.451 | 82.12 | 80.14 | 87.72 | 35.55 | 0.312 | 32.10 |
| GLM-5.3-Flash [39] | 68.02 | 0.111 | 0.344 | 81.86 | 83.70 | 70.71 | 43.09 | 0.288 | 39.09 |
| *Specialized VLMs (OCR-2.0)* | | | | | | | | | |
| Unlimited-OCR [37] | 61.89 | 0.167 | 0.605 | 73.71 | 73.51 | 60.76 | 53.19 | 0.204 | 31.58 |
| *Specialized VLMs (OCR-3.0)* | | | | | | | | | |
| **Unlimited-OCR 2 (Ours)** | **84.59** | **0.052** | **0.100** | **89.39** | **87.78** | **89.54** | **72.91** | **0.077** | **60.10** |

Table 2 | Performance comparison on MarathonDocBench. Unlimited OCR is not reported on MarathonDocBench due to its 40-page input limitation.

| Method | Overall↑ | Text$^{Edit}$ ↓ | T-Text$^{Edit}$ ↓ | Formula$^{CDM}$ ↑ | Table$^{TEDS}$ ↑ | T-Table$^{TEDS}$ ↑ | Figure$^{F1}$ ↑ | ReadOrder$^{Edit}$ ↓ | Heading$^{TEDS}$ ↑ |
|---|---|---|---|---|---|---|---|---|---|
| *General VLMs* | | | | | | | | | |
| Gemini-3.1-Pro Preview [9] | 26.97 | 0.516 | 0.764 | 30.44 | 28.48 | 21.71 | 18.94 | 0.654 | 9.68 |
| Qwen3.8-Max [25] | 22.93 | 0.600 | 0.804 | 18.82 | 15.20 | 13.71 | 33.39 | 0.733 | 16.09 |
| Kimi-K3 [24] | 33.74 | 0.475 | 0.664 | 40.50 | 27.80 | 26.30 | 21.38 | 0.611 | 28.90 |
| Minimax-M3 [16] | 17.97 | 0.659 | 0.818 | 14.20 | 17.79 | 14.84 | 3.67 | 0.763 | 17.16 |
| Gemini-3.8-Flash [10] | 28.61 | 0.519 | 0.658 | 26.35 | 23.40 | 19.56 | 18.62 | 0.687 | 27.34 |
| Qwen-3.8-Flash-Next [25] | 27.83 | 0.544 | 0.753 | 28.31 | 25.38 | 21.04 | 24.07 | 0.679 | 21.50 |
| DeepSeek-4.1-Flash [32] | 12.34 | 0.740 | 0.933 | 10.45 | 12.74 | 12.59 | 6.95 | 0.843 | 7.62 |
| GLM-5.3-Flash [39] | 7.51 | 0.860 | 0.936 | 4.50 | 7.12 | 6.65 | 7.10 | 0.923 | 6.55 |
| *Specialized VLMs (OCR-3.0)* | | | | | | | | | |
| **Unlimited-OCR 2 (Ours)** | **77.06** | **0.104** | **0.123** | **87.93** | **74.81** | **71.11** | **71.79** | **0.147** | **48.20** |

### 5.2. Main Results

As visually summarized in Figure 6 and further detailed in Table 1 and Table 2, Unlimited OCR 2 takes a significant step toward the OCR 3.0 era, advancing document parsing along three key dimensions. First, it establishes the practical feasibility of hundred-page end-to-end parsing, achieving a score of 77.06 on MarathonDocBench and demonstrating its ability to handle documents at a scale beyond conventional multi-page settings. Second, it maintains robust parsing performance as document length increases. Through end-to-end modeling with SR-SWA, Unlimited OCR 2 filters redundant visual tokens while preserving relevant cross-page information, achieving the highest Overall scores and state-of-the-art performance on both cross-page text and table metrics, T-Text and T-Table, respectively, from dozens of pages on MPDocBench to hundreds of pages on MarathonDocBench. Third, it establishes a new state-of-the-art in long-document parsing, outperforming leading closed-source models and top-tier open-source MLLMs across these benchmarks. Furthermore, Unlimited OCR 2 retains strong single-page recognition capabilities, achieving an Overall score of 94.02% on OmniDocBench in Table 3, confirming that its sparse attention design preserves fine-grained local recognition.